Solvency Field Theory V10.3: Unified Maintenance Closure

The Corrected Critical Variable, the Four-Regime History, and the Robustness Basin

A standalone extension paper building on V10, V10.1, and V10.2

April 2026

Purpose. V10 named the unification of the shatter and the engine crash as an open theorem-level target. The missing object was not more rhetoric but a corrected critical variable and a dynamical bridge between the two epochs. This paper preserves that closure cluster. Theorem 1a repairs the unified solvency ratio and puts the early and late transitions into one sign-consistent maintenance law. Theorem 1b then shows that the resulting history is not a two-event analogy but a four-regime maintenance sequence

$$\text{shared} \to \text{independent} \to \text{active} \to \text{passive},$$

with the shatter and the engine crash as the two channel-critical reorganizations inside one globally continuous envelope. The paper also preserves the first robustness basin and the first sensitivity ranking of the load-bearing directions.

Abstract

This extension preserves the first successful unification program for the two forced transitions of Solvency Field Theory. First, we repair the critical variable itself. The unified solvency ratio must use the same orientation already present in the forced-transition scaffold and the shatter sector,

$$\Lambda_X(t) = \frac{H(t)I(t)}{C_X(t)} = \frac{D(t)}{C_X(t)},$$

not its inverse. The shatter and the engine crash are then two channel-critical crossings of the same maintenance law under different dominant servicing channels. Second, we show that a three-channel history is incomplete: it leaves a post-shatter insolvent gap. The missing object is an independent or seed-structure bridge channel. Once that channel is included, the effective history becomes four-regime,

$$\text{shared} \to \text{independent} \to \text{active} \to \text{passive},$$

and the dominant maintenance envelope remains solvent throughout the successful realizations. Third, we preserve the first numerical closure support for this picture. A pinned effective sweep yields 489 accepted models out of 972 tested (acceptance fraction 0.503). A softer local-release sweep still preserves 1836 accepted models out of 8748 tested (acceptance fraction 0.210). Finally, a one-parameter marginal sensitivity scan identifies the first load-bearing ranking of the four channels: shared-channel weakening is the most fragile direction, the independent bridge has finite-window

slack, the active channel has moderate room, and the passive channel is the most tolerant direction in the scanned box. The honest status after this paper is therefore: Theorem 1a closes the corrected scaffold; Theorem 1b is supported at the effective closure level with a real robustness basin; and the remaining frontier is no longer existence, but first-principles derivation of the channel histories and the structural bridge.

Contents

1 Why V10.3 exists

V10 named the unification of the shatter and the engine crash as an open theorem target. The idea was already visible: both events were being described as forced transitions of the maintenance process. What was still missing was the kind of object a theorem actually needs. The sign convention was not yet locked. The denominator was not yet given one stable physical meaning. The bridge between the early and late transitions was still mostly narrative. And the sequence of maintenance regimes between them had not yet been made explicit.

This paper exists because that situation changed.

The first change was conceptual but decisive. The unified solvency ratio in the early fragment was written as capacity over demand. That was incompatible with the orientation already used in the generic forced-transition scaffold and in the shatter sector. The repair is simple,

$$\Lambda_X = \frac{D}{C_X},$$

but it matters because it keeps all stability language aligned:

$$\Lambda < 1 \text{ solvent}, \qquad \Lambda = 1 \text{ critical}, \qquad \Lambda > 1 \text{ insolvent}.$$

The second change was dynamical. A three-channel toy history can place the shatter in the intended thermal window and the late active crash near the phantom-crossing neighborhood, but it cannot carry the universe cleanly from one to the other. It leaves a post-shatter insolvent interval in the dominant envelope. That failure is informative. It identifies the missing object: an independent or seed-structure bridge channel.

The third change was numerical. Once the bridge channel is added, the four-regime picture is not just possible; it occupies a real effective basin. That is the point where a theorem cluster deserves preservation.

This paper therefore does for Theorem 1 what V10.2 did for the projection cluster. It does not try to rewrite the whole framework. It preserves the closure steps that have now earned their own extension.

2 The argument in one line

The argument of this paper is short enough to state directly:

> The shatter and the engine crash are not two unrelated stories. They are the two channel-critical reorganizations of one maintenance process once the unified solvency ratio is written correctly and the missing post-shatter bridge channel is included.

The sections that follow unpack that sentence in the same discipline-first style the framework has been learning to use:

- Section 3 states Theorem 1a and repairs the critical variable.
- Section 4 shows why the three-channel history fails and why the independent bridge is required.
- Section 5 preserves Theorem 1b: the first successful four-regime effective closure.
- Section 6 preserves the robustness basin.
- Section 7 preserves the sensitivity ranking.
- Section 8 states exactly what is now closed and what remains frontier work.

3 Theorem 1a: corrected unified critical variable

3.1 The sign convention had to be repaired

The early unified-solvency fragment identified the right physical target — a single ratio linking the shatter and the engine crash — but wrote it in the wrong orientation. Inside the established forced-transition scaffold, the control variable is demand over capacity, not capacity over demand. The shatter sector already used the same convention in local form. So the first theorem repair is not decorative. It is load-bearing.

Define the universal ledger demand by

$$D(t) \equiv H(t)I(t) = \dot{S}_{\text{total}}(t),$$

and define the channel-specific maintenance ratio by

$$\Lambda_X(t) = \frac{D(t)}{C_X(t)} = \frac{H(t)I(t)}{C_X(t)},$$

where C_X is the servicing capacity of the channel currently trying to carry the maintenance burden.

The interpretation is fixed once and then used everywhere:

$$\Lambda_X < 1 \quad \text{solvent / viable,}$$

$$\Lambda_X = 1 \quad \text{critical,}$$

$$\Lambda_X > 1 \quad \text{insolvent / forced reorganization.}$$

This is the first closure step because without it the early and late sectors cannot even share one theorem statement cleanly.

3.2 The two epoch-specific crossings

With the repaired orientation, the early-universe shared regime is governed by

$$\Lambda_{\text{sh}}(T) = \frac{H(T)I(T)}{\Gamma_{\text{shared}}(T)}.$$

The shatter occurs when

$$\Gamma_{\text{shared}}(T_c) = H(T_c)I(T_c).$$

The late-universe active regime is governed by

$$\Lambda_{\text{act}}(z) = \frac{H(z)I(z)}{\dot{S}_{\text{active}}(z)}, \qquad \dot{S}_{\text{active}}(z) = \dot{S}_{\text{BH}}(z) + \dot{S}_{\text{SF}}(z).$$

The engine crash occurs when

$$\dot{S}_{\text{active}}(z_{\text{crash}}) = H(z_{\text{crash}})I(z_{\text{crash}}).$$

The point is not that the numerators are identical. The point is that they are both capacities measured against the same denominator.

3.3 One denominator, not two

The dangerous conceptual point in the original fragment was denominator drift. The theorem cannot survive if HI means one thing at the shatter and another thing at the crash. The clean statement is:

$$D(t) = H(t)I(t)$$

is always the required irreversible maintenance throughput of the realized universe at epoch t.

At early times, that throughput is experienced as the burden the shared thermal regime must collectively carry. At late times, it is the burden the active astrophysical engines would need to carry to remain dominant. Those are two physical descriptions of one ledger object, not two different demands that happen to share a formula.

3.4 Theorem 1a

Theorem 1a. *There exists a single SFT critical variable*

$$\Lambda_X(t) = \frac{H(t)I(t)}{C_X(t)}$$

such that the shatter and the engine crash are its first two channel-critical crossings under different dominant servicing channels, with the early event carried by the shared channel and the late event carried by the active channel.

Theorem 1a does not yet prove the full history between those two crossings. It repairs the scaffold that makes the full theorem possible.

4 Why the three-channel history fails

The natural first attempt after Theorem 1a is to model a three-channel history:

$$\text{shared} \to \text{active} \to \text{passive}.$$

That attempt is useful because it fails in exactly the right way.

A three-channel toy can do several things correctly:

- it can place the shatter in the intended thermal window,
- it can place the late active crash near the observed phantom-crossing neighborhood,
- and it can let the passive structural regime rise later.

But it still leaves a long post-shatter interval in which the dominant maintenance envelope is insolvent. In other words, the model knows how to begin and how to end, but not how to cross the long middle.

That failure identifies the missing object immediately. The universe does not jump straight from shared thermal servicing to a mature active astrophysical era. There must be an intermediate regime in which independent ticking exists, organized structure can begin to grow, and the later active engines become possible.

So the missing channel is

$$C_{\text{ind}}(z),$$

the independent or seed-structure bridge channel.

That is not a patch added after the fact. It is the dynamical version of the four-era story the framework was already trying to tell.

5 Theorem 1b: four-regime effective closure

5.1 The four regimes

Once the independent bridge channel is included, the effective maintenance history becomes

$$\text{shared} \to \text{independent} \to \text{active} \to \text{passive}.$$

The role of each regime is different:

Regime	Carrier	Role in the theorem
Shared	C_{sh}	Collective thermal servicing before the shatter
Independent	C_{ind}	Post-shatter bridge, seed-structure era, enables later active servicing
Active	C_{act}	Main interaction era, baryonic engines dominate, later crash
Passive	C_{pass}	Structural or geometric handoff era after active depletion

This is the first point where the theorem stops looking like two disconnected events and starts looking like one maintenance history with memory.

5.2 Effective channel forms

Theorem 1b is not yet a first-principles derivation of all four channel histories. It is an effective closure program. The point is to use the mildest channel shapes consistent with the intended physical logic.

The shared channel is controlled by a shatter ratio of the form

$$\Lambda_{\mathrm{sh}}(z) = \lambda_{\mathrm{lo}} + (\lambda_{\mathrm{hi}} - \lambda_{\mathrm{lo}}) \frac{1}{1 + \left(\frac{1+z}{1+z_{\mathrm{sh}}}\right)^p},$$

with

$$C_{\mathrm{sh}}/D = 1/\Lambda_{\mathrm{sh}}.$$

This separates location from stiffness cleanly: z_{sh} fixes the crossing location, while p controls the sharpness.

The structural bridge is encoded in a monotone fraction

$$f_{\mathrm{struct}}(z),$$

which is still imposed rather than derived, but now appears as an explicit theorem variable instead of a purely narrative bridge.

The independent channel turns on at the shatter and declines only gradually once the active engines can take over. The active channel is modeled as a structure-gated rise followed by a late crash. The passive channel is a late-rising structure-supported reservoir.

The important point is not the exact algebraic form of each toy function. The important point is that once the independent bridge is included, the successful histories stop opening dominant insolvent gaps.

5.3 Baseline effective result

In the successful baseline realization, the dominant history is exactly

$$\text{shared} \to \text{independent} \to \text{active} \to \text{passive}.$$

The key numerical landmarks are:

- shatter crossing at $z \approx 3301.15$, corresponding to $T_{\text{sh}} \approx 9000$ K,
- independent-to-active handoff at $z \approx 1.88$,
- active late crossing at $z \approx 0.4645$,
- passive crossing at $z \approx 0.4899$,
- active-to-passive handoff at $z \approx 0.625$,
- and no dominant insolvent interval.

This is the central effective-closure result.

5.4 Channel criticality versus global solvency

The four-regime model clarifies something important. A specific channel can become critical when

$$C_X/D = 1,$$

without forcing the global maintenance envelope to become insolvent. The shatter is the critical failure of the shared channel. The engine crash is the critical failure of the active channel. But the maintenance process as a whole can remain globally solvent because another channel is already available to take over.

That is a better theorem structure than "the universe becomes insolvent twice." What the model actually shows is:

> The universe undergoes two channel-critical reorganizations inside one globally continuous maintenance history.

5.5 Theorem 1b

Theorem 1b (effective existence form). *There exists a four-channel effective realization of the corrected SFT critical-condition scaffold in which the maintenance process passes through the ordered regime history*

$$\text{shared} \to \text{independent} \to \text{active} \to \text{passive},$$

the shatter and the engine crash remain the two channel-critical transitions, and the dominant maintenance envelope remains solvent throughout the successful realizations.

This is not the final first-principles theorem. It is the first successful effective closure of the four-regime picture.

6 Robustness basin

6.1 Pinned-basin scan

The first robustness test kept the successful four-channel scaffold relatively pinned and asked whether the intended history occupied more than a single tuned point. It did.

Out of 972 total models, 489 were accepted, for an acceptance fraction of

$$0.503.$$

That result matters because it shows the four-regime picture is not a knife-edge construction when the scaffold is tested in the neighborhood where it was first discovered.

6.2 Soft-release scan

The next test loosened the successful scaffold locally rather than freeing everything at once. That is the correct question. The failed hard-release sweep had shown that unconstrained amplitudes can erase the active era entirely and reopen dominant insolvent gaps. The soft-release scan instead asks whether the successful theorem picture survives local slack.

It does.

Out of 8748 total models, 1836 were accepted, for an acceptance fraction of

$$0.210.$$

The accepted models preserve the exact regime order

$$\text{shared} \to \text{independent} \to \text{active} \to \text{passive},$$

with no dominant insolvent interval.

Robustness pass	Accepted / Total	Acceptance fraction
Pinned basin	489 / 972	0.503
Soft release	1836 / 8748	0.210

That is already enough to move Theorem 1b from exploratory worksheet to serious preserved result.

7 Marginal sensitivity ranking

7.1 Why the marginal scan matters

Once the soft-release basin exists, the next question is no longer whether the theorem picture exists. The next question is which directions are actually load-bearing. A one-parameter-at-a-time scan around the best accepted soft-release model answers that.

The slack variables are

$$\epsilon_{\mathrm{sh}}, \qquad \epsilon_{\mathrm{ind}}, \qquad \epsilon_{\mathrm{act}}, \qquad \epsilon_{\mathrm{pass}}.$$

7.2 Shared channel

The shared-channel slack is strongly one-sided.

- Accepted count: 16 / 31
- Accepted range:
$$\epsilon_{\mathrm{sh}} \in [0.00,\ 0.30]$$
- Negative-side failure mode: dominant insolvent gap

Interpretation: weakening the shared channel destabilizes the pre-shatter handoff almost immediately. The shared regime must remain strong enough to carry the maintenance burden right up to the shatter.

7.3 Independent bridge channel

The independent bridge channel is finite-window load-bearing.

- Accepted count: 12 / 31
- Accepted range:
$$\epsilon_{\mathrm{ind}} \in [0.00,\ 0.22]$$
- Negative-side failure mode: dominant insolvent gap
- Positive-side failure mode: sequence break

Interpretation: if the bridge is weakened, the post-shatter envelope opens a gap. If it is strengthened too much, it persists too long and obstructs the clean handoff into the active regime.

7.4 Active channel

The active channel has moderate but finite slack.

- Accepted count: 21 / 31

- Accepted range:

$$\epsilon_{\text{act}} \in [-0.18,\, 0.22]$$

- Negative-side failure mode: sequence break
- Positive-side failure mode: active-window failure

Interpretation: the active era is robust, but not arbitrarily so. If the active channel is weakened too much, it cannot establish the intended regime history. If it is strengthened too much, the late crash leaves the intended critical neighborhood.

7.5 Passive channel

The passive channel is the most tolerant direction in the scanned box.

- Accepted count: 31 / 31
- Accepted range:

$$\epsilon_{\text{pass}} \in [-0.30,\, 0.30]$$

Interpretation: the passive structural regime is the least fragile part of the current effective scaffold.

7.6 Sensitivity ranking

The marginal scan therefore gives the first clean ranking of the load-bearing directions:

1. most fragile: shared-channel weakening,
2. next most fragile: independent bridge-channel slack,
3. moderately constrained: active-channel amplitude,
4. most tolerant: passive-channel amplitude.

The theorem picture is no longer only an existence story. It now distinguishes which channels are structurally constrained and which channels possess real local slack.

8 Updated status ledger

Theorem 1 now has a cleaner internal status split.

Item	Status	Meaning
Theorem 1a	Closed scaffold	Corrected critical variable, unified sign convention, one denominator meaning
Theorem 1b	Effective closure support	Four-regime history, robustness basin, sensitivity ranking
Channel microphysics	Open	First-principles derivation of the effective channel histories
Structural bridge law	Open	Derivation of $f_{\text{struct}}(z)$ or equivalent monotonicity law
Full uniqueness / exact-two proof	Open	Stronger proof that only the shatter and engine crash are channel-critical reorganizations before heat death

The right language now is not "fully proved" but "effectively closed enough to preserve."

9 What remains frontier work

Three real theorem tasks remain.

9.1 First-principles channel derivations

The current channel shapes are effective and intentionally mild. They preserve the intended logic, but they are not yet microscopic or semi-microscopic derivations.

9.2 Structural bridge derivation

The bridge variable $f_{\text{struct}}(z)$ is no longer just narrative, but it is still imposed rather than derived. A stronger closure step will need either an actual evolution law or a theorem-level monotonicity argument derived from the framework itself.

9.3 Exact-two statement

The current theorem cluster supports two true channel-critical reorganizations of the successful maintenance history. It does not yet prove the strongest possible statement that these are the only such reorganizations before heat death. That is a later theorem frontier.

A Full Theorem 1a note (preserved form)

Theorem 1a. There exists a single SFT critical variable

$$\Lambda_X(t) = \frac{H(t)I(t)}{C_X(t)}$$

such that the early-universe shatter and the late-universe engine crash are its first two channel-critical crossings under two different dominant servicing channels. The denominator

$$D(t) = H(t)I(t)$$

has one fixed meaning: required irreversible maintenance throughput of the realized universe at epoch t. The two crossings are causally linked because the first transition enables the growth of organized structure, and that structure in turn enables the active servicing channel whose later depletion drives the second transition.

What Theorem 1a closes is the corrected sign convention, the unified denominator interpretation, and the equation-level candidate object. What it does not yet close is the full intermediate history. That is why Theorem 1b exists.

B Four-regime effective definitions (preserved form)

The four effective channels used in the worksheet are summarized schematically by

$$C_{\mathrm{sh}}/D = \Lambda_{\mathrm{sh}}^{-1},$$

$$C_{\mathrm{ind}} \propto \text{turn-on at shatter} \times \text{slow decline after } z_{\mathrm{drop}},$$

$$C_{\mathrm{act}} \propto f_{\mathrm{struct}}^{\alpha_{\mathrm{struct}}} \times \text{turn-on} \times \text{crash},$$

$$C_{\mathrm{pass}} \propto f_{\mathrm{struct}}^{\beta_{\mathrm{struct}}} \times \text{late rise}.$$

The point of preserving these forms is not that they are final microphysics. The point is that they are the first effective channel set that can carry the theorem fragment's intended causal chain without opening a dominant insolvent interval.

C Preserved numerical landmarks

The following values are the key preserved landmarks of the successful theorem cluster.

Quantity	Preserved value
Best shatter temperature anchor	$T_{\mathrm{sh}} \approx 9000$ K
Best shatter crossing	$z \approx 3301.15$
Baseline independent-to-active handoff	$z \approx 1.88$
Best active late crossing	$z \approx 0.4645$
Best passive crossing	$z \approx 0.4899$
Baseline active-to-passive handoff	$z \approx 0.625$
Pinned-basin acceptance fraction	$489/972 = 0.503$
Soft-release acceptance fraction	$1836/8748 = 0.210$

D Bottom line

The theorem fragment's causal chain is no longer just a plausible story.

The unified maintenance history

$$\text{shared} \to \text{independent} \to \text{active} \to \text{passive}$$

survives over a real basin of effective parameter space, and that basin is not uniformly fragile. The shared and independent sectors are the most load-bearing parts of the scaffold; the active sector has finite room; the passive sector is the most tolerant.

So the current honest status is this: Theorem 1a closes the corrected scaffold, Theorem 1b has strong effective closure support, and the next frontier is no longer existence but derivation.